

- 1- Prepare Kubeadm env (EC2s on AWS)
- 2- Docker (Manual Installation)
- 3- Nodejs app
- 4- Prometheus & Grafana(Manual Installation)

```
ubuntu@masternode:~$ kubectl get ns
NAME          STATUS   AGE
default       Active   3h22m
kube-flannel   Active   3h21m
kube-node-lease Active   3h22m
kube-public    Active   3h22m
kube-system    Active   3h22m
monitoring     Active   3h17m
```

Open the required ports using kubectl expose commands.

```
ubuntu@masternode:~$ kubectl get svc -n monitoring
NAME                                TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)                                AGE
alertmanager-operated              ClusterIP    None             <none>            9093/TCP,9094/TCP,9094/UDP            3h17m
prometheus-grafana                 ClusterIP    10.105.94.63     <none>            80/TCP                                3h17m
prometheus-grafana-ext             NodePort     10.111.141.235   <none>            80:30403/TCP                          3h7m
prometheus-kube-prometheus-alertmanager ClusterIP    10.103.161.117   <none>            9093/TCP,8080/TCP                    3h17m
prometheus-kube-prometheus-alertmanager-ext NodePort     10.102.253.183   <none>            9093:32677/TCP,8080:30930/TCP        171m
prometheus-kube-prometheus-operator ClusterIP    10.107.215.193   <none>            443/TCP                                3h17m
prometheus-kube-prometheus-prometheus ClusterIP    10.99.167.195     <none>            9090/TCP,8080/TCP                    3h17m
prometheus-kube-prometheus-prometheus-ext NodePort     10.108.204.107   <none>            9090:30238/TCP,8080:31353/TCP        3h11m
prometheus-kube-state-metrics      ClusterIP    10.98.129.203     <none>            8080/TCP                                3h17m
prometheus-operated                ClusterIP    None             <none>            9090/TCP                                3h17m
prometheus-prometheus-node-exporter ClusterIP    10.98.156.62     <none>            9100/TCP                                3h17m
```

Access the Prometheus and Grafana UIs using IP:NodePort. (Grafana password: prom--operator)

```
ubuntu@masternode:~$ mkdir nodejs
ubuntu@masternode:~$ cd nodejs/
ubuntu@masternode:~/nodejs$
```

```
ubuntu@masternode:~$ vim index.js
```

```
const express = require('express'); // web framework to create a lightweight server
const client = require('prom-client'); // prometheus library to collect metrics

const app = express();
const port = 3000; // open port 3000

// Create a Registry to register the metrics
const register = new client.Registry();

// Add a default label which is added to all metrics
register.setDefaultLabels({
  app: 'nodejs_system_app'
});

// Enable the collection of default metrics
client.collectDefaultMetrics({ register }); // Memory - Cpu - event loop lag - gc state

// Define a custom metric for total HTTP requests to the root path
const rootHttpRequestCounter = new client.Counter({
  name: 'http_requests_root_total',
  help: 'Total number of HTTP requests to the root path',
});

// Register the custom metric
register.registerMetric(rootHttpRequestCounter);

// Middleware to count every request to the root path
app.use((req, res, next) => {
  if (req.path === '/') {
    rootHttpRequestCounter.inc();
  }
  next();
});

// Define a route for Prometheus to scrape
app.get('/metrics', async (req, res) => { // open the endpoint /metrics that store information
  res.set('Content-Type', register.contentType);
  res.end(await register.metrics());
});

// Define the root route
app.get('/', (req, res) => {
  res.send('Hello From Node.js app Test');
});

// Start the server
app.listen(port, () => {
  console.log(`Example app listening at http://localhost:${port}`);
});
```

```
ubuntu@masternode:~$ vim Dockerfile
```

```
FROM node:lts
WORKDIR /usr/src/app
COPY . .
RUN npm install express prom-client
EXPOSE 3000
CMD ["node", "index.js"]
```

```
ubuntu@masternode:~/nodejs$ docker build -t nodejs-app-prom .
```

```
[+] Building 5.1s (9/9) FINISHED
=> [internal] load build definition from Dockerfile
=> => transferring dockerfile: 156B
=> [internal] load metadata for docker.io/library/node:lts
=> [internal] load .dockerignore
=> => transferring context: 2B
=> [1/4] FROM docker.io/library/node:lts@sha256:e515259afd26f60db74957c62203c93d45760f2ba864d94accfa2edfc1ac17cf
=> [internal] load build context
=> => transferring context: 184B
=> CACHED [2/4] WORKDIR /usr/src/app
=> [3/4] COPY .
=> [4/4] RUN npm install express prom-client
=> exporting to image
=> => exporting layers
=> => writing image sha256:21391370b13e5047165456ae10376cce0ea6f1161ece3b2f4d73639e5bbb09dd
=> => naming to docker.io/library/nodejs-app-prom
```

```
ubuntu@masternode:~/nodejs$ docker images
REPOSITORY          TAG          IMAGE ID          CREATED           SIZE
nodejs-app-prom     latest       21391370b13e     About a minute ago 1.14GB
```


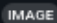
```
ubuntu@masternode:~/nodejs$ docker tag nodejs-app-prom:latest aliaceofali/nodejs-app-prom:v1
```

```
ubuntu@masternode:~/nodejs$ docker images
REPOSITORY          TAG          IMAGE ID          CREATED           SIZE
aliaceofali/nodejs-app-prom  v1          21391370b13e     About a minute ago 1.14GB ## check for tag
nodejs-app-prom     latest       21391370b13e     About a minute ago 1.14GB
```

```
docker login
Login Succeeded
## verification
```

```
ubuntu@masternode:~/nodejs$ docker push aliaceofali/nodejs-app-prom:v1
The push refers to repository [docker.io/aliaceofali/nodejs-app-prom]
f515febcb2f4: Pushed
7617c43270a7: Pushed
8735979c33fc: Pushed
e5b77f05d0e5: Mounted from library/node
bd8201ae8f5d: Mounted from library/node
16af5fc601df: Mounted from library/node
2f10455f8757: Mounted from library/node
1b90aaab596: Mounted from library/node
a9b908787288: Mounted from library/node
20c8aca21338: Mounted from library/node
175a19836175: Mounted from library/node
v1: digest: sha256:6b99fe3af88641a2eb4d0de2f8bc95122127efa37d93a188ca4ebc915880afe1 size: 2630
```

Check your image on your DockerHub account.

Name	Last Pushed 	Contains	Visibility	Scout
aliaceofali/nodejs-app-prom	less than a minute ago 		Public	Inactive

Deployment Steps

- **Create a deployment.**
- **Expose it and check access using the NodePort.**

```
ubuntu@masternode:~/nodejs$ kubectl get deploy
No resources found in default namespace.
```

```
ubuntu@masternode:~/nodejs$ kubectl get po
No resources found in default namespace.
```

```
ubuntu@masternode:~$ vim nodejs-app.yaml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nodejs-app
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nodejs
  template:
    metadata:
      labels:
        app: nodejs
    spec:
      containers:
        - name: nodejs
          image: aliaceofali/nodejs-app-prom:v1
          ports:
            - containerPort: 3000
```

```
ubuntu@masternode:~/nodejs$ kubectl apply -f nodejs-app.yaml
deployment.apps/nodejs-app created
```

```
ubuntu@masternode:~/nodejs$ kubectl get deploy
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
nodejs-app    3/3     3            3           27s
```

```
ubuntu@masternode:~/nodejs$ kubectl get rs
NAME                                DESIRED   CURRENT   READY   AGE
nodejs-app-8549c8b8db              3         3         3       29s
```

```
ubuntu@masternode:~/nodejs$ kubectl get po
NAME                                READY   STATUS    RESTARTS   AGE
nodejs-app-8549c8b8db-cjj7v        1/1     Running   0          23s
nodejs-app-8549c8b8db-twmrn        1/1     Running   0          23s
nodejs-app-8549c8b8db-zgk5q        1/1     Running   0          23s
```

```
apiVersion: v1
kind: Service
metadata:
  name: nodejs-svc
  labels:
    app: nodejs
  annotations:
    prometheus.io/scrape: "true" # Enable scraping for Prometheus
spec:
  type: NodePort
  selector:
    app: nodejs
  ports:
    - port: 3000
      targetPort: 3000
      name: http-nodejs-app
```

```
ubuntu@masternode:~/nodejs$ kubectl apply -f nodejs-svc.yaml
service/nodejs-svc created
```

```
kubectl get svc
ubuntu@masternode:~/nodejs$ kubectl get svc
NAME          TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE
kubernetes    ClusterIP   10.96.0.1    <none>        443/TCP          3h45m
nodejs-svc    NodePort    10.97.14.36  <none>        3000:31557/TCP   11s
```

```
ip : 31557(NodePort)
```

Hello From Node.js app Test

Enable Prometheus to Monitor the Application (Not Just the Pods)

```
ubuntu@masternode:~/nodejs$ kubectl get po
NAME                                READY   STATUS    RESTARTS   AGE
nodejs-app-8549c8b8db-cjj7v        1/1     Running   0           3m7s
nodejs-app-8549c8b8db-twmrn        1/1     Running   0           3m7s
nodejs-app-8549c8b8db-zgk5q        1/1     Running   0           3m7s
```

In the Prometheus UI, execute the query:
`container_cpu_usage_seconds_total`

- Use the Pod ID to search for results.
- You'll find how much CPU the application is using.
- Prometheus **did not** scrape metrics from the application – only from the pod itself.

```
container_cpu_usage_seconds_total(cpu="total", endpoint="https-metrics", id="/kubepods.slice/kubepods-besteffort.slice/kubepods-
besteffort-poddf1d95a6_a39b_4536_a85a_437eb1add8ef.slice/cri-
containerd-3a601598f37386d5bb4412dc4e9f71c0af235e577f0ba863c1510b515ff56a1e.scope", image="k8s.gcr.io/pause:3.5",
instance="172.31.38.31:10250", job="kubelet", metrics_path="/metrics/cadvisor",                                0.030135
name="3a601598f37386d5bb4412dc4e9f71c0af235e577f0ba863c1510b515ff56a1e", namespace="default", node="node1", pod="nodejs-
app-8549c8b8db-cjj7v", service="prometheus-kube-prometheus-kubelet")
```

So let's tell Prometheus about the application by adding it as a **new target**.

To do this, we need to create a **ServiceMonitor** component in the **Prometheus namespace**.

```
ubuntu@masternode:~/nodejs$ kubectl get ns
NAME                STATUS   AGE
default             Active   3h49m
kube-flannel        Active   3h49m
kube-node-lease     Active   3h49m
kube-public         Active   3h49m
kube-system         Active   3h49m
monitoring          Active   3h44m
```

```
ubuntu@masternode:~/nodejs$ kubectl apply -f nodejs-monitor.yaml

apiVersion: monitoring.coreos.com/v1
kind: ServiceMonitor
metadata:
  name: nodejs-monitor
  namespace: monitoring # Ensure this matches your monitoring namespace
  labels:
    release: prometheus # Adjust if you use a different labels ( depends on system )
spec:
  selector:
    matchLabels:
      app: nodejs # Match the label used in your nodejs-app.yaml
  namespaceSelector:
    matchNames:
      - default # where your nodejs-svc is deployed
  endpoints:
    - port: http-nodejs-app # Match the name of the port in your nodejs-svc.yaml
      path: /metrics # Path to scrape metrics from
```

```
ubuntu@masternode:~/nodejs$ kubectl apply -f nodejs-monitor.yaml
servicemonitor.monitoring.coreos.com/nodejs-monitor created
```

Open Prometheus and check the **Targets** page to see if the new application endpoint is added.

serviceMonitor/monitoring/nodejs-monitor/0				3 / 3 up		^
Endpoint	Labels	Last scrape	State			
http://10.244.1.9:3000/metrics	container="nodejs"	2.242s ago	12ms	UP		

```
ubuntu@masternode:~/nodejs$ kubectl get po -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP            NODE   NOMINATED NODE   READINESS GATES
nodejs-app-8549c8b8db-cjj7v        1/1     Running   0           11m   10.244.1.9    node1   <none>            <none>
nodejs-app-8549c8b8db-twmrn        1/1     Running   0           11m   10.244.2.8    node2   <none>            <none>
nodejs-app-8549c8b8db-zgk5q        1/1     Running   0           11m   10.244.2.9    node2   <none>            <none>
```

You'll find the endpoint listed under **Targets**, along with its IP. Now Prometheus can monitor both the pod and the application inside it.

Add a New Target Using ServiceMonitor

- The prometheus-operator will automatically discover and start scraping the new target.

The screenshot displays the Prometheus Targets page with three discovered targets. Each target is represented by a card showing its URL, labels, and status.

Target URL	Labels	Scrape Time	Duration	Status
http://10.244.1.9:3000/metrics	container="nodejs" endpoint="http-nodejs-app" instance="10.244.1.9:3000" job="nodejs-svc" namespace="default" pod="nodejs-app-8549c8b8db-cjj7v" service="nodejs-svc"	2.242s ago	12ms	UP
http://10.244.2.8:3000/metrics	container="nodejs" endpoint="http-nodejs-app" instance="10.244.2.8:3000" job="nodejs-svc" namespace="default" pod="nodejs-app-8549c8b8db-twmrn" service="nodejs-svc"	1.75s ago	7ms	UP
http://10.244.2.9:3000/metrics	container="nodejs" endpoint="http-nodejs-app" instance="10.244.2.9:3000" job="nodejs-svc" namespace="default" pod="nodejs-app-8549c8b8db-zgk5q" service="nodejs-svc"	3.247s ago	11ms	UP

Create a Rule

- Let's now create an alerting rule.
- You can check configured rules in the **UI under Alerts/Rules**.
- Configure the rule to send alerts to **Slack**.
- Create a Bash script to send many requests to the application, simulating **stress or load**.
- Observe the behavior in the UI – even in networking namespaces if applicable.

```
apiVersion: monitoring.coreos.com/v1
kind: PrometheusRule
metadata:
  name: nodejs-alerts
  namespace: monitoring # Ensure this matches your monitoring namespace
  labels:
    app: kube-prometheus-stack
    release: prometheus # it must match the label used in your Prometheus default configuration
spec:
  groups:
    - name: nodejs.alert # Name of the alert group
      rules:
        - alert: HighRequestRate_NodeJS
          expr: rate(http_requests_root_total[5m]) > 10 # expression means high request rate
          for: 0m
          labels:
            app: nodejs
            namespace: monitoring # Ensure this matches your monitoring namespace
          annotations:
            summary: "High request rate detected in Node.js application"
            description: "The Node.js application is receiving a high request rate ({{ $value }})"

ubuntu@masternode:~/nodejs$ kubectl apply -f nodejs-rule.yaml
prometheusrule.monitoring.coreos.com/nodejs-alerts created
```

The screenshot displays the Prometheus Alerting interface for a rule named 'nodejs.alert'. At the top, the rule's file path is shown: '/etc/prometheus/rules/prometheus-prometheus-kube-prometheus-prometheus-rulefiles-0/monitoring-nodejs-alerts-360badd3-1527-4a83-8071-401cd75ce0a1.yaml'. Below this, status indicators show 'last run 10.256s ago', 'took 0ms', and 'every 30s'. The main section shows the alert name 'HighRequestRate_NodeJS' with a bell icon, a refresh icon, and a green 'OK' button. The expression 'rate(http_requests_total[5m]) > 10' is displayed in a dark box. Below the expression, the labels 'app="nodejs"' and 'namespace="monitoring"' are shown. The 'description' field contains 'The Node.js application is receiving a high request rate ({{ \$value }})' and the 'summary' field contains 'High request rate detected in Node.js application'.

```

ubuntu@masternode:~/nodejs$ kubectl get secret -n monitoring
NAME                                     TYPE      DATA      AGE
alertmanager-prometheus-kube-prometheus-alertmanager      Opaque    1          3h56m
alertmanager-prometheus-kube-prometheus-alertmanager-cluster-tls-config Opaque    1          3h56m
alertmanager-prometheus-kube-prometheus-alertmanager-generated Opaque    1          3h56m
alertmanager-prometheus-kube-prometheus-alertmanager-tls-assets-0 Opaque    0          3h56m
alertmanager-prometheus-kube-prometheus-alertmanager-web-config Opaque    1          3h56m
prometheus-grafana                                         Opaque    3          3h56m
prometheus-kube-prometheus-admission                     Opaque    3          3h56m
prometheus-prometheus-kube-prometheus-prometheus          Opaque    1          3h56m
prometheus-prometheus-kube-prometheus-prometheus-thanos-prometheus-http-client-file Opaque    1          3h56m
prometheus-prometheus-kube-prometheus-prometheus-tls-assets-0 Opaque    1          3h56m
prometheus-prometheus-kube-prometheus-prometheus-web-config Opaque    1          3h56m
sh.helm.release.v1.prometheus.v1                         helm.sh/release.v1 1          3h56m
slack-secret                                               Opaque    1          3h6m

```

```

ubuntu@masternode:~/nodejs$ kubectl describe secret slack-secret -n monitoring
Name:      slack-secret
Namespace: monitoring
Labels:    <none>
Annotations: <none>

Type: Opaque

Data
====
webhook: 81 bytes

```

```

ubuntu@masternode:~/nodejs$ vim nodejs-alert-manager.yaml
apiVersion: monitoring.coreos.com/v1alpha1
kind: AlertmanagerConfig
metadata:
  name: nodejs-alert-manager
  namespace: monitoring # Ensure this matches your monitoring namespace
spec:
  route:
    receiver: 'nodejs-slack'
    repeatInterval: 30m
    routes:
      - matchers:
          - name: alert-https
            value: HighRequestRate_NodeJS # Match the alert name you created in nodejs-rule.yaml
            repeatInterval: 10m
      receivers:
        - name: 'nodejs-slack'
          slackConfigs:
            - apiURL:
                key: webhook
                name: slack-secret # Ensure this matches your secret name
                channel: '#https_requests' # Adjust to your Slack channel
                sendResolved: true

```

```

ubuntu@masternode:~/nodejs$ kubectl apply -f nodejs-alert-manager.yaml
alertmanagerconfig.monitoring.coreos.com/nodejs-alert-manager created

```

```

ubuntu@masternode:~/nodejs$ kubectl get alertmanagerconfig -n monitoring
NAME                                AGE
nodejs-alert-manager                35s

```

```

ubuntu@masternode:~/nodejs$ kubectl get svc
NAME      TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
kubernetes ClusterIP  10.96.0.1      <none>         443/TCP          4h12m
nodejs-svc NodePort    10.97.14.36   <none>         3000:31557/TCP   27m

```

```

ubuntu@masternode:~/nodejs$ vim forcerequests.sh

#!/bin/bash

send_requests() {
  while true; do
    curl -sS http://13.61.174.226:31557/ > /dev/null
    echo "Request sent"
    sleep 0.0667
  done
}

for ((i=1; i<=15; i++)); do
  send_requests &
  pids[$i]=$!
done

trap 'echo "Exiting..."; kill ${pids[*]}; exit' INT

wait

ubuntu@masternode:~/nodejs$ sudo chmod +x forcerequests.sh
ubuntu@masternode:~/nodejs$ ./forcerequests.sh
Request sent
Request sent
Request sent

```

Add Visualization in Grafana

- Take the Prometheus query expression and add it to Grafana.
- Watch the graph showing your application's metrics in real-time.
- This helps you visualize request traffic and performance data.

FIRING (1)

HighRequestRate_NodeJS

FIRING (3)

```
rate(http_requests_root_total[5m]) > 10
```

app="nodejs"namespace="monitoring"

description

The Node.js application is receiving a high request rate ({{{ \$value }}})

summary

High request rate detected in Node.js application

Alert labels

State

Active Since

Value

alername="HighRequestRate_NodeJS"app="nodejs"container="nodejs"

endpoint="http-nodejs-app"instance="10.244.2.9:3000"job="nodejs-svc"

namespace="monitoring"pod="nodejs-app-8549c8b8db-zgk5q"service="nodejs-svc"

FIRING

3.422s

40.41111111111111

description

The Node.js application is receiving a high request rate (40.41111111111111)

summary

High request rate detected in Node.js application

alername="HighRequestRate_NodeJS"app="nodejs"container="nodejs"

endpoint="http-nodejs-app"instance="10.244.1.9:3000"job="nodejs-svc"

namespace="monitoring"pod="nodejs-app-8549c8b8db-cjj7v"service="nodejs-svc"

FIRING

3.422s

40.42962962962963

description

The Node.js application is receiving a high request rate (40.42962962962963)

summary

High request rate detected in Node.js application

alername="HighRequestRate_NodeJS"app="nodejs"container="nodejs"

endpoint="http-nodejs-app"instance="10.244.2.8:3000"job="nodejs-svc"

namespace="monitoring"pod="nodejs-app-8549c8b8db-twrmn"service="nodejs-svc"

FIRING

3.422s

40.262962962962966

description

The Node.js application is receiving a high request rate (40.262962962962966)


summary

High request rate detected in Node.js application

monitoring/master-alert-manager/Sendithere-slack namespace="monitoring" 5 alerts

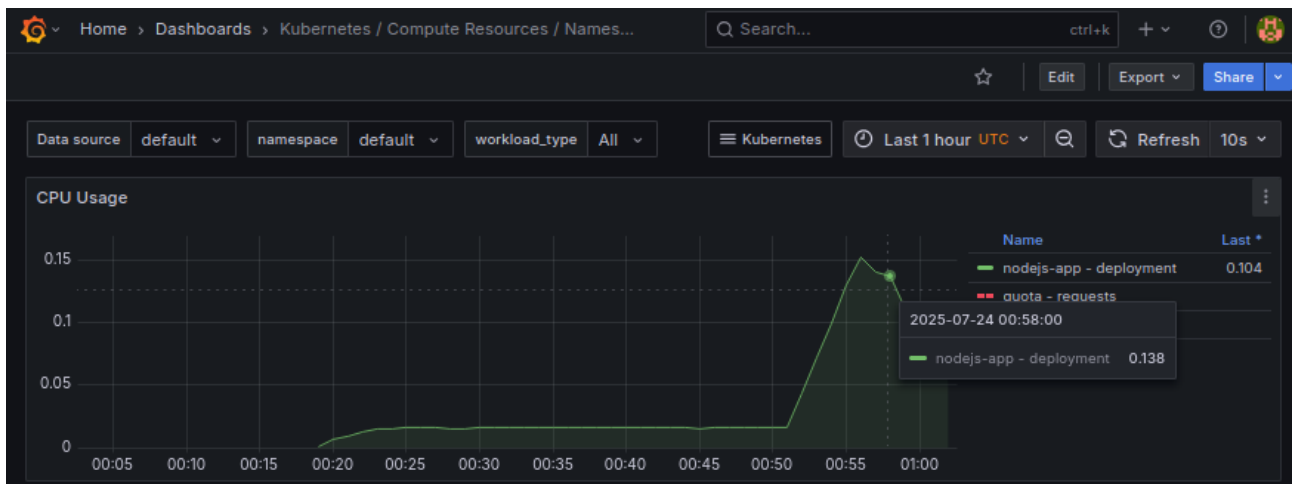
2025-07-24T00:59:22.974Z + Info Source Silence Link

alertname="HighRequestRate_NodeJS" app="nodejs" container="nodejs" endpoint="http-nodejs-app" instance="10.244.1.9:3000" job="nodejs-svc" pod="nodejs-app-8549c8b8db-cjj7v" prometheus="monitoring/prometheus-kube-prometheus-prometheus" service="nodejs-svc"

 **cpu-usage** APP 3:58 AM

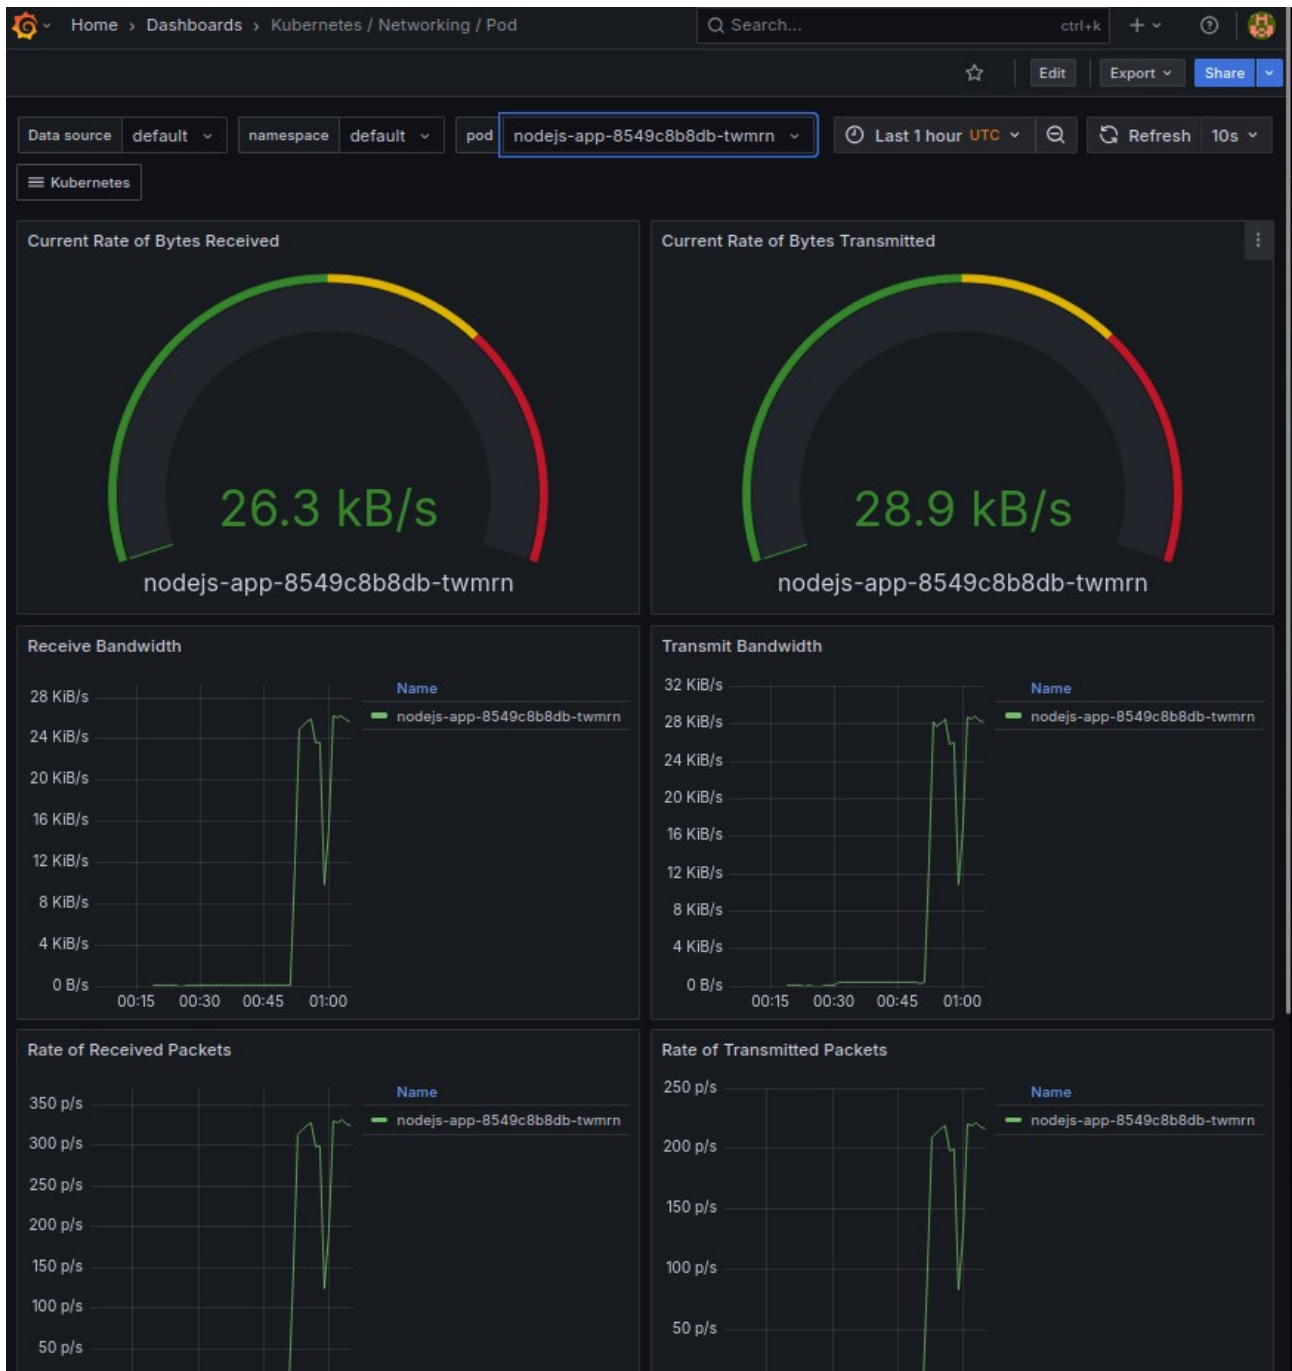
[FIRING:1] monitoring (ConfigReloaderSidecarErrors config-reloader reloader-web 10.244.1.8:8080 prometheus-kube-prometheus-alertmanager alertmanager-prometheus-kube-prometheus-alertmanager-0 monitoring/prometheus-kube-prometheus-prometheus prometheus-kube-prometheus-alertmanager warning)

React Reply



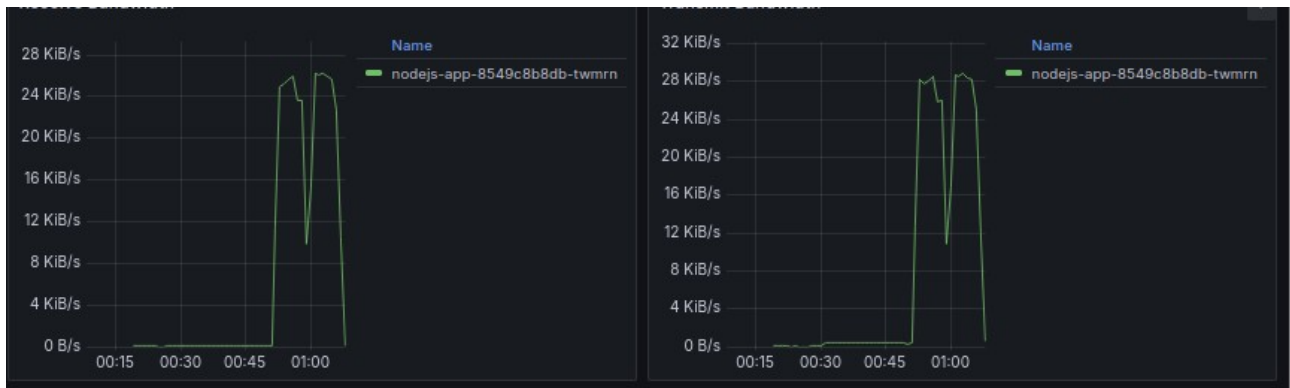
let's stop the script





Stop the Bash Script

- Once the stress test is done, stop the script.
- Then check the **Grafana graphs** again to observe how the application metrics respond.



1.